

BDA Expt 7

HBASE

Hadoop Database



Google File System

To solve
distributed
storage

MapReduce

To solve
distributed
computing

Google File System

MapReduce

**Apache developed
open source versions
of these technologies**

Google File System




HDFS

MapReduce



MapReduce

Hadoop



**A file system to manage
the storage of data**

**A framework to process data
across multiple servers**

Hadoop is a big data processing
framework

Hadoop is **not** a database!

The Importance of Databases

What Kind of Data Do Organizations Store?



Order Management

An e-commerce site stores order information



Payroll

A company stores employee payroll details



Accounts

A bank stores account and transaction information



Requirements of a Database

Structured: Rows and columns

Random access: Update one row at a time

Low latency: Very fast read/write/update operations

ACID compliant: Ensure data integrity

What Are ACID Properties?

Atomicity

Consistency

Isolation

Durability

Atomicity

Transactions on a database should be **all-or-nothing**

Atomicity

Transferring Money




Both withdrawal and deposit should occur or none at all!

Consistency

Database updates
should not violate
any **constraints**

Consistency

Enrolling Students



**Every student should
have a unique student id**

Isolation

Concurrent operations
on the database should
appear as though they
were applied in some
sequence

Isolation

Granularity of Updates



**Can employee address be
updated at the same time
as employee salary?**

Durability

Once changes have
been made to the
data they are
permanent

Durability

Safety of Data



**In case of power
loss, crashes, errors**

Limitations of Hadoop

Unfortunately, Hadoop makes a very poor database

Limitations of Hadoop



Unstructured data



No random access



High latency



Not ACID compliant

Limitations of Hadoop

Basic structure exists
for some file types

CSV files

XML files

JSON files

Hadoop enforces no
constraints on these



Unstructured data

Limitations of Hadoop

Cannot create, access
and modify individual
records in a file

MapReduce parses
entire files to extract



No random access

Limitations of Hadoop

Not suited for real-time processing where a user waits for data to be retrieved

Batch processing with long running jobs



High latency

Limitations of Hadoop

HDFS is a file storage system and provides no guarantees for data integrity




Not ACID compliant

How Did Google Solve This for Search?



Google published a paper
on Bigtable a distributed
storage system for
structured data

How Did Google Solve This for Search?



HBase is a distributed database management system which runs on top of Hadoop

HBase



HBase

Distributed: Stores data in HDFS

Scalable: Capacity directly proportional to number of nodes in the cluster

Fault tolerant: Piggybacks on Hadoop


Structured: A loose data structure

Low latency: Real-time access using row based indices called row keys

Random access: Row keys allow access updates to one record

Somewhat ACID compliant: Some transactions will have ACID properties

Properties of HBase



Columnar store



Denormalized storage




Only CRUD operations



ACID at the row level

Columnar Store



Columnar store

Columnar Store

Id	To	Type	Content
1	mike	offer	Offer on mobiles
2	john	sale	Redmi sale
3	jill	order	Order delivered
4	megan	sale	Clothes sale



Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Offer on mobiles
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	megan
4	Type	sale
4	Content	Clothes sale

Advantages of a Columnar Store

Sparse tables: No wastage of space when storing sparse data


Dynamic attributes: Update attributes dynamically without changing storage structure

Sparse Tables

Id	To	Type	Content	Expiry	Order Status
1	mike	offer	Offer on mobiles	2345689070	
2	john	sale	Redmi sale		
3	jill	order	Order delivered		Delivered
4	megan	sale	Clothes sale	2456123989	

And empty cells when data is not applicable to certain rows

These cells still occupy space.




Columnar store

Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Offer on
1	Expiry	2345689070
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	megan
4	Type	sale
4	Content	Clothes sale
4	Expiry	2456123989

Dynamically add new attributes as rows in this table

No wastage of space with empty cells!

Properties of HBase



Columnar store



Denormalized storage



Only CRUD operations



ACID at the row level

Denormalized Storage

Id	Name	Function	Grade
1	Emily	Finance	6

Id	Subordinate Id
1	2
1	3

Id	City	Zip Code
1	Palo Alto	94305
2	Seattle	98101

Normalization

- ▣ Normalization optimizes storage,
- ▣ But storage is cheap in distributed system.
- ▣ We have to optimize no. Of disk seeks.

Denormalized Storage

Id	Name	Function	Grade
1	Emily	Finance	6
2	John	Finance	3
3	Ben	Finance	4

Id	Subordinate Id
1	2
1	3



Id	Name	Function	Grade	Subordinates <ARRAY>
1	Emily	Finance	6	
2	John	Finance	3	
3	Ben	Finance	4	

Denormalized Storage

Id	Name	Function	Grade
1	Emily	Finance	6
2	John	Finance	3
3	Ben	Finance	4

Id	City	Zip Code
1	Palo Alto	94305
2	Seattle	98101




Id	Name	Function	Grade	Subordinates	Address
1	Emily	Finance	6	<ARRAY>	<STRUCT>
2	John	Finance	3		
3	Ben	Finance	4		

Store everything related to an employee in the same table

Read a single record to get all details about an employee in one read operation

Properties of HBase



Columnar store



Denormalized storage



Only CRUD operations



ACID at the row level



Only CRUD operations

Traditional Databases and SQL

Joins: Combining information across tables using keys

Group By: Grouping and aggregating data for the groups

Order By: Sorting rows by a certain column



Only CRUD operations

**HBase does
not support
SQL**

NoSQL



Only CRUD operations

Only a limited set of operations
are allowed in HBase

Create

Read

Update

Delete

CRUD




Only CRUD operations

Id	Name	Function	Grade	Subordinates	Address

This is why all details
need to be **self
contained** in one row

Properties of HBase



Columnar store



Denormalized storage



Only CRUD operations



ACID at the row level



ACID at the row level

Updates to a single row are atomic

All columns in a row are updated or none are



ACID at the row level

Updates to multiple rows are **not** atomic

Even if the update is on the same column in multiple rows

Traditional RDBMS vs. HBase

Traditional RDBMS

Data arranged in rows and columns

Supports SQL

Complex queries such as grouping, aggregates, joins etc

Normalized storage to minimize redundancy and optimize space

ACID compliant

HBase

Data arranged in a column-wise manner

NoSQL database

Only basic operations such as create, read, update and delete

Denormalized storage to minimize disk seeks

ACID compliant at the row level

How Is Data Laid out in HBase?

Traditional database

Id	To	Type	Content
1	mike	offer	Offer on mobiles
2	john	sale	Redmi sale
3	jill	order	Order delivered
4	megan	sale	Clothes sale

This is a 2-dimensional data model

HBase has a **4-dimensional**
data model

4 Dimensions



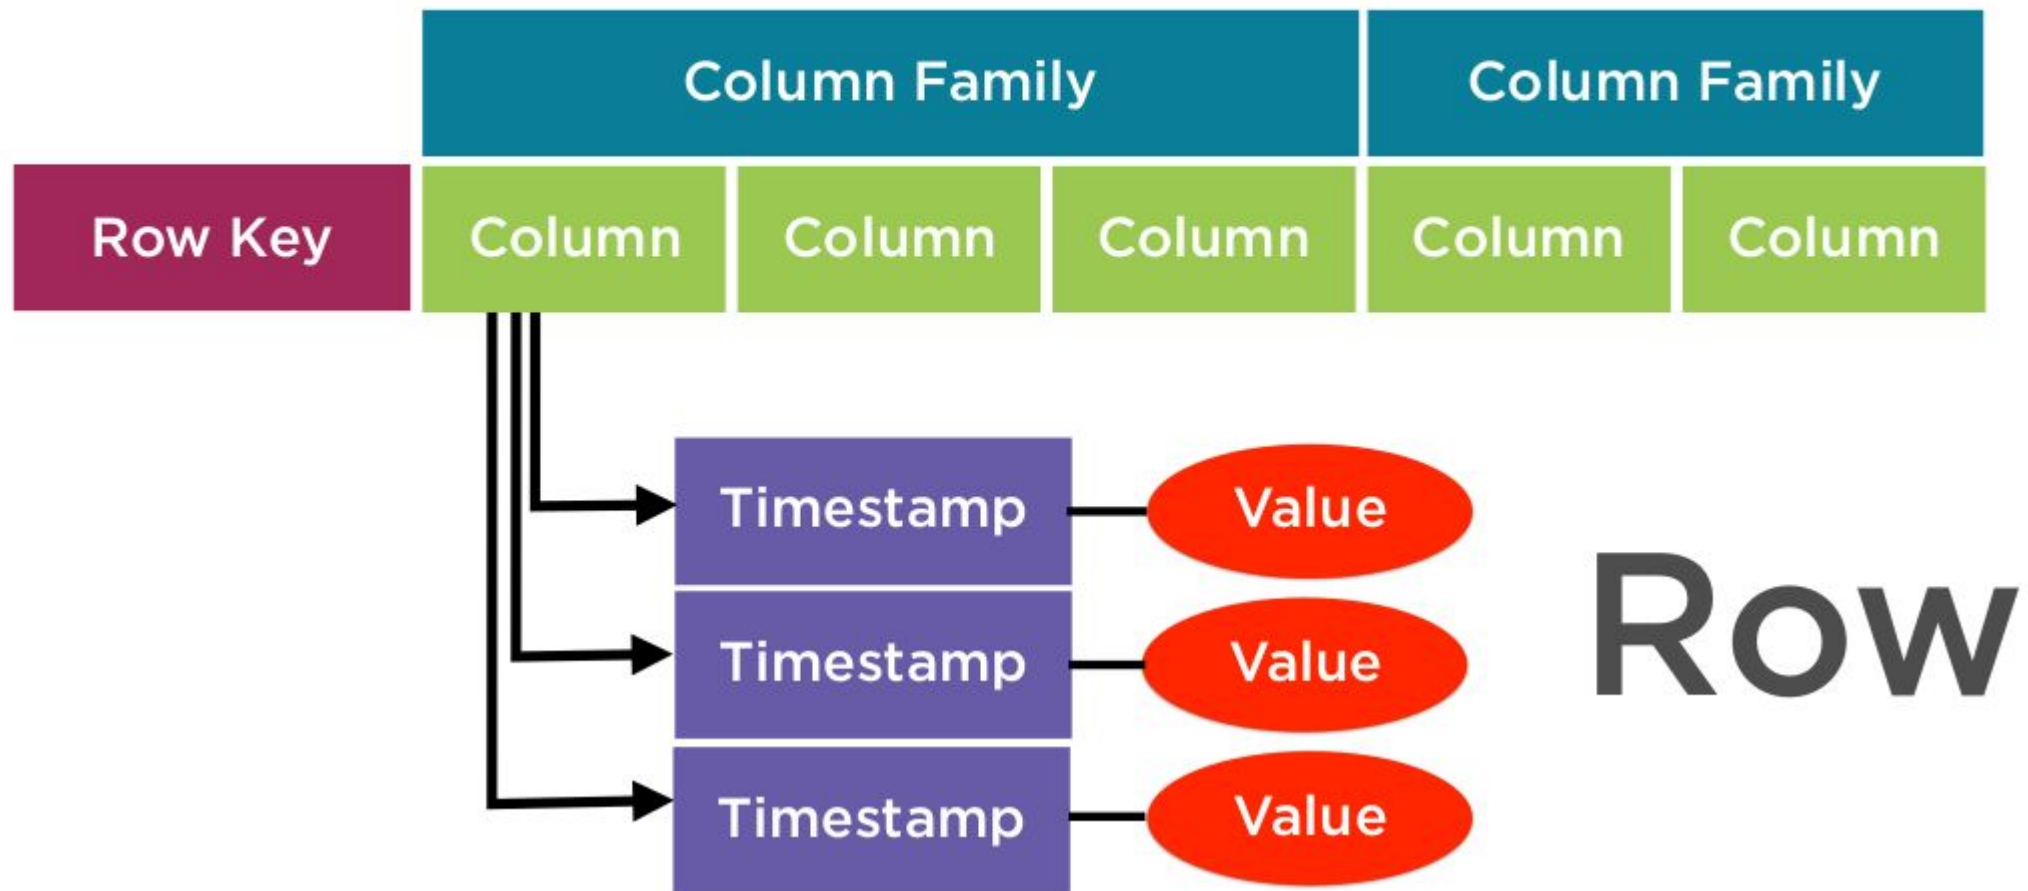
Row Key

Column Family

Column

Timestamp

4-dimensional Data Model



Census Data Layout in HBase

Some ID	Personal			Professional	
	name	gender	marital_status	employed	field

Notification Data

Id	To	Type	Content
1	mike	offer	Offer on mobiles
2	john	sale	Redmi sale
3	jill	order	Order delivered
4	megan	sale	Clothes sale



Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Offer on mobiles
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	megan
4	Type	sale
4	Content	Clothes sale

Row Key



Row Key

Uniquely identifies a row

Can be primitives, structures, arrays

Represented internally as a byte array

Sorted in ascending order

Id	To	Type	Content
1	mike	offer	Offer on mobiles
2	john	sale	Redmi sale
3	jill	order	Order delivered
4	megan	sale	Clothes sale



Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Offer on mobiles
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	megan
4	Type	sale
4	Content	Clothes sale

Column Family

A teal square graphic with the text "Column Family" centered inside it.

Column Family

All rows have the same set of column families

Each column family is stored in a separate data file

Set up at schema definition time

Can have different columns for each row



Column

Columns are units within a column family

New columns can be added on the fly

ColumnFamily: ColumnName =
Work:Department



Timestamp

Used as the version number for the values stored in a column

The value for any version can be accessed

Hbase shell commands

Census Data Layout in HBase

Some ID	Personal			Professional	
	name	gender	marital_status	employed	field

- put 'census', 1, 'personal:name', 'Mike Jones'
- put 'census', 1, 'personal:marital_status',
'unmarried'
- put 'census', 1, 'personal:gender', 'male'
- put 'census', 1, 'professional:employed', 'yes'
- put 'census', 1, 'professional:education_level',
'high school'
- put 'census', 1, 'professional:field', 'construction'

- put 'census', 3, 'personal:name', 'Jill Tang'
- put 'census', 3, 'personal:marital_status', 'married'
- put 'census', 3, 'personal:spouse', 'Jim Tang'
- put 'census', 3, 'professional:education_level',
'post-grad'
- put 'census', 3, 'personal:gender', 'female'
- put 'census', 3, 'personal:name', 'Ben'

```
hbase(main):003:0> list
TABLE
Employees
Sales
notifications
3 row(s) in 0.0480 seconds
```

```
hbase(main):004:0> create 'census', 'personal', 'professional'
```

```
hbase(main):008:0> count 'census'
0 row(s) in 0.0510 seconds
```

```
hbase(main):007:0> describe 'census'
```

```
Table census is ENABLED
```

```
census
```

```
COLUMN FAMILIES DESCRIPTION
```

```
{NAME => 'personal', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '1', BLOOMFILTER_CACHED => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

```
{NAME => 'professional', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '1', BLOOMFILTER_CACHED => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

```
2 row(s) in 0.1170 seconds
```

Insert and update data using the HBase shell

- Add rows, each row represents data specific to one person
- Edit cells in a row

```
hbase(main):001:0> put 'census' 1, 'personal:name', 'Mike Jones'
```

The row key can be any data structure

```
hbase(main):001:0> put 'census', 1, 'personal:name', 'Mike Jones'
```

Insert data one cell at a time

The column family prefix for every column qualifier

```
put 'census', 1, 'personal:marital_status', 'unmarried'|
```

Another column inserted for the same row

```
hbase(main):003:0> scan 'census'
ROW          COLUMN+CELL
 1          column=personal:marital_status, timestamp=1479390670583, value=unmarried
 1          column=personal:name, timestamp=1479390660627, value=Mike Jones
1 row(s) in 0.0520 seconds
```

The timestamp serves as a version number for the value in the cell

The "put" command can be used to **update** values in cells as well

By default HBase always retrieves the value with the latest timestamp

- Syntax of put is same for insertion and updation.