



IIT KHARAGPUR
IIT GANDHINAGAR



NPTEL ONLINE
CERTIFICATION COURSES

Scaling Data Science

Lecture 6: Introduction to Hashing

Anirban Dasgupta

Computer Science and Engineering
IIT GANDHINAGAR



IIT Gandhinagar
Indian Institute of
Technology Gandhinagar

Outline

- Outline:
 - Hash tables and hash functions
 - Universal hashing
 - Chaining
 - Multiplicative hashing



NPTEL



Querying



Present?



Naïve algorithm: linear in dataset size

Hash Table

- Elements come from universe U , but we need to store only n items, $n < |U|$
- Hash table
 - array of size m
 - Hash function $h: U \rightarrow \{0, 1, \dots, m-1\}$
- We typically use $m \ll |U|$ as well as $m < n$
 - Collisions happen when $x \neq y$, but $h(x) = h(y)$

Hash functions

- In theory, we design for worst-case behaviour of data
 - Need to choose hash function “randomly”
- Hash family $H = \{h_1, h_2, \dots\}$
 - When creating hash table, a **single** function $h \in H$ is chosen randomly
 - We then analyse the expected query time
- However...



Hash functions

- In theory, we design for worst-case behaviour of data
 - Need to choose hash function “randomly”
- Hash family $H = \{h_1, h_2, \dots\}$
 - When creating hash table, a **single** function $h \in H$ is chosen randomly
 - We then analyse the expected query time
- Since the algo has to carry around the “description” of the hash function, it needs $\log(|H|)$ bits of storage
 - $|H|$ cannot too big, in particular, it cannot be the set $[m]^U$, all possible functions

Hash family

- We need to create small hash families H such that choosing from it gives a function with “good behaviour”

NPTEL



Hash family

- We need to create small hash families H such that choosing from it gives a function with “good behaviour”
- **Uniform:** $\Pr_{h \in H} [h(x) = i] = \frac{1}{m}$ for all x and i



Hash family


- We need to create small hash families H such that choosing from it gives a function with “good behaviour”
- Uniform: $\Pr_{h \in H}[h(x) = i] = \frac{1}{m}$ for all x and i
 - Not enough
- Universal: $\Pr_h[h(x) = h(y)] = \frac{1}{m}$ for all $x \neq y$

Hash family

- We need to create small hash families H such that choosing from it gives a function with “good behaviour”
- Uniform: $\Pr_{h \in H}[h(x) = i] = \frac{1}{m}$ for all x and i
 - Not enough
- Universal: $\Pr_h[h(x) = h(y)] = \frac{1}{m}$ for all $x \neq y$
- Near Universal: $\Pr_h[h(x) = h(y)] \leq \frac{2}{m}$ for all $x \neq y$

Chaining

- When collisions happen, we store elements using a linked list from that location



Chaining

- When collisions happen, we store elements using a linked list from that location
- $l(x)$ = length of chain at position $h(x)$
- Expected time to query $x = O(1 + E_h[l(x)])$
 - Same for insert and delete



Analyzing chaining

- Need to bound $E_h[l(x)]$
- For $x \neq y$, define $C_{xy} = \begin{cases} 1 & \text{if } h(x) = h(y) \\ 0 & \text{else} \end{cases}$



Analyzing chaining

- Need to bound $E_h[l(x)]$
- For $x \neq y$, define $C_{xy} = \begin{cases} 1 & \text{if } h(x) = h(y) \\ 0 & \text{else} \end{cases}$
- $E_h[l(x)] = E_h[\sum_y C_{xy}]$

Analyzing chaining: universal hashing

- Need to bound $E_h[l(x)]$
- For $x \neq y$, define $C_{xy} = \begin{cases} 1 & \text{if } h(x) = h(y) \\ 0 & \text{else} \end{cases}$
- $E_h[l(x)] = E_h[\sum_y C_{xy}] = \sum_y \Pr[h(x) = h(y)] = \frac{n}{m}$
universal

Multiplicative hashing

- How to design small + universal hash family?
- Prime multiplicative hashing:
 - Fix a prime number $p > |U|$
 - $H = \{ h_a(x) = (ax \bmod p) \bmod m, a \in \{1, \dots, p-1\} \}$
 - Choosing a hash function is same as choosing $a \in \{1, \dots, p-1\}$

Multiplicative hashing

- $H = \{ h_a(x) = (ax \bmod p) \bmod m, a \in \{1, \dots, p-1\} \}$
- This family satisfies $\Pr_h[h(x) = h(y)] \leq \frac{1}{m}$
- Intuition: $h_a(x) - h_a(y) = (a(x - y) \bmod p) \bmod m$
- There are at most $\frac{p-1}{m}$ values in $\{1, \dots, p-1\}$ that are divisible by m

Multiplicative hashing

- $H = \{ h_a(x) = (ax \bmod p) \bmod m, a \in \{1, \dots, p-1\} \}$
- This family satisfies $\Pr_h[h(x) = h(y)] \leq \frac{1}{m}$
- **Intuition:** $h_a(x) - h_a(y) = (a(x - y) \bmod p) \bmod m$
- There are at most $\frac{p-1}{m}$ values in $\{1, \dots, p-1\}$ that are divisible by m
- What is the probability of choosing a such that $(a(x - y) \bmod p)$ is one of these numbers?

A property of prime numbers

WLOG $x - y \in [1, p - 1]$

Property: For every $t, z \in [1, p - 1]$ there exists unique $a \in [1, p - 1]$ such that $az \bmod p = t$

This would imply that probability of choosing collision-causing a

$$\leq \frac{p-1}{m} \times \frac{1}{p-1} = \frac{1}{m}$$

A property of prime numbers

WLOG $x - y \in [1, p - 1]$

Property: For every $t, z \in [1, p - 1]$ there exists unique $a \in [1, p - 1]$ such that $az \bmod p = t$

By contradiction. If not, then $\exists a, b \in [1, p - 1]$ such that $(a - b)z \bmod p = 0$.

But this cannot be as p is prime.

k-wise universal

- For any distinct (x_1, \dots, x_k) and any (not necessarily distinct) (y_1, \dots, y_k) ,

$$\Pr[h(x_1) = y_1 \wedge \dots \wedge h(x_k) = y_k] = m^{-k}$$

- Needs only $O(k \log n)$ bit of storage

NPTEL



Summary

- Hashing
 - Simple and versatile
 - Main issue is design of good hash functions, much researched area
 - (near) universality guarantees small chain sizes
 - Other alternatives to chaining exist, e.g. open addressing, cuckoo hashing



References:

- Primary reference for this lecture
 - Algorithms and models of computation by Jeff Erickson:
<http://jeffe.cs.illinois.edu/teaching/algorithms/>
- Others
 - Algorithms, by Cormen, Leiserson and Rivest
 - Randomized Algorithms by Mitzenmacher and Upfal.

Thank You!!



IIT Gandhinagar
Indian Institute of
Technology Gandhinagar



NPTEL ONLINE
CERTIFICATION COURSES

Anirban Dasgupta
Computer Science and Engg.



IIT KHARAGPUR
IIT GANDHINAGAR



NPTEL ONLINE
CERTIFICATION COURSES

Scalable Data Science

Lecture 7: Bloom Filters

Anirban Dasgupta

Computer Science and Engineering


IIT GANDHINAGAR





IIT Gandhinagar

Indian Institute of
Technology Gandhinagar


Querying



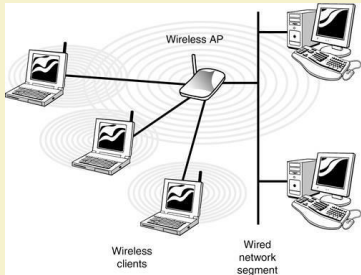
ISBN present in collection?



IP seen by switch?



10.0.21.102



Solutions

- Universe U , but need to store a set of n items, $n \ll |U|$
- Hash table of size m :
 - Space $O(n \log |U|)$
 - Query time $O\left(\frac{n}{m}\right)$



Solutions

- Universe U , but need to store a set of n items, $n \ll |U|$
- Hash table of size m :
 - Space $O(n \log |U|)$
 - Query time $O\left(\frac{n}{m}\right)$
- Bit array of size $|U|$
 - Space = $|U|$
 - Query time $O(1)$



Querying, Monte Carlo style


- In hash table construction, we used random hash functions
 - we never return incorrect answer
 - query time is a random variable
 - These are Las Vegas algorithms
- In Monte-Carlo randomized algorithms, we are allowed to return incorrect answers with (small) probability, say, δ



Bloom filter


[Bloom, 1970]

- A bit-array B , $|B| = m$
- k hash functions, h_1, h_2, \dots, h_k , each $h_i \in U \rightarrow [m]$



Bloom filter

- A bit-array B , $|B| = m$
- k hash functions, h_1, h_2, \dots, h_k , each $h_i \in U \rightarrow [m]$



Operations

- *Initialize*(B)
 - for $i \in \{1, \dots, m\}$, $B[i] = 0$
- *Insert* (B, x)
 - for $i \in \{1, \dots, k\}$, $B[h_i(x)] = 1$
- *Lookup* (B, x)
 - If $\bigwedge_{i \in \{1, \dots, k\}} B[h_i(x)]$, return PRESENT, else ABSENT



Bloom Filter


- If the element x has been added to the Bloom filter, then $Lookup(B, x)$ always return PRESENT

NPTEL



Bloom Filter

- If the element x has been added to the Bloom filter, then $Lookup(B, x)$ always return PRESENT
- If x has not been added to the filter before?
 - $Lookup$ sometimes still return PRESENT



Designing Bloom Filter

- Want to minimize the probability that we return a false positive
- Parameters $m = |B|$ and $k =$ number of hash functions
- $k = 1 \Rightarrow$ normal bit-array
- What is effect of changing k ?



Effect of number of hash functions

- Increasing k
 - Possibly makes it harder for false positives to happen in *Lookup* because of $\bigwedge_{i \in \{1, \dots, k\}} B[h_i(x)]$
 - But also increases the number of filled up positions
- We can analyse to find out an “optimal k ”



False positive analysis

- $m = |B|$, n elements inserted
- If x has not been inserted, what is the probability that $Lookup(B, x)$ returns PRESENT?

NPTEL



False positive analysis

- $m = |B|$, n elements inserted
- If x has not been inserted, what is the probability that $Lookup(B, x)$ returns PRESENT?
- Assume $\{h_1, h_2, \dots, h_k\}$ are independent and $\Pr[h_i(\cdot) = j] = \frac{1}{m}$ for all positions j
- $\Pr[h_i(x) = 0] = \left(1 - \frac{1}{m}\right)^{kn} \approx e^{-kn/m}$

False positive analysis

- The expected number of zero bits $\approx me^{-kn/m}$ w.h.p.
- $\Pr[\text{Lookup}(B, x) = \text{PRESENT}] = (1 - e^{-kn/m})^k$
- Can we choose k to minimize this probability


Choosing number of hash functions

- $p = e^{-kn/m}$

- Log (False Positive) =

$$\log(1 - p)^k = k \log(1 - p) = -\frac{m}{n} \log(p) \log(1 - p)$$

Minimized at $p = \frac{1}{2}$, i.e. $k = m \log(2)/n$



Bloom filter design

- This “optimal” choice gives false positive = $2^{-m \log(2)/n}$
- If we want a false positive rate of δ , set $m = \left\lceil \frac{\log\left(\frac{1}{\delta}\right)n}{\log^2(2)} \right\rceil$

Example: If we want 1% FPR, we need 7 hash functions and total $10n$ bits

Applications

- Widespread applications whenever small false positives are tolerable
- Used by browsers
 - to decide whether an URL is potentially malicious: a BF is used in browser, and positives are actually checked with the server.
- Databases e.g. BigTable, HBase, Cassandra, Postgrepsql use BF to avoid disk lookups for non-existent rows/columns
- Bitcoin for wallet synchronization....

Handling deletions

- Chief drawback is that BF does not allow deletions
- Counting Bloom Filter [Fan et al 00]
 - Every entry in BF is a small counter rather than a single bit
 - $Insert(x)$ increments all counters for $\{h_i(x)\}$ by 1
 - $Delete(x)$ decrements all $\{h_i(x)\}$ by 1
 - maintains 4 bits per counter
 - False negatives can happen, but only with low probability



Other Extensions

- Many recent work on Bloom filters
 - Can we do with less hashing?
 - Can BFs be compressed (needed for distributed systems)
 - Are there better structures that use less space, less randomness and less memory lookups?



References:

- Primary reference for this lecture
 - Survey on Bloom Filter, Broder and Mitzenmacher 2005,
<https://www.eecs.harvard.edu/~michaelm/postscripts/im2005b.pdf>
 - <http://www.firatatagun.com/blog/2016/09/25/bloom-filters-explanation-use-cases-and-examples/>
- Others
 - Randomized Algorithms by Mitzenmacher and Upfal.



Thank You!!



IIT Gandhinagar
Indian Institute of
Technology Gandhinagar



NPTEL ONLINE
CERTIFICATION COURSES

Anirban Dasgupta
Computer Science and Engg.



IIT KHARAGPUR
IIT GANDHINAGAR



NPTEL ONLINE
CERTIFICATION COURSES

Scalable Data Science

Lecture 8: Streaming model, counting distinct elements

Anirban Dasgupta

Computer Science and Engineering
IIT GANDHINAGAR



IIT Gandhinagar
Indian Institute of
Technology Gandhinagar

Large Data

- Data is massive, growing faster than our ability to store or index
- Predicted growth of data @ 1.7Mb/person/second [Forbes]
- Scientific data:
 - Large Hadron Collider
 - Gravitational wave detector
 - Personalized genome sequences

Handling velocity + volume

- Can we process data without explicitly storing all of it in memory? E.g. in a network switch,
 - which IPs have most packets passing through a switch
 - has traffic pattern changed overnight?



Handling velocity + volume

- Can we process data without explicitly storing all of it in memory? E.g. in a network switch,
 - which IPs have most packets passing through a switch
 - has traffic pattern changed overnight?
- We have to give up on exact answer, and rely on...
 - approximation: return answer close to truth
 - randomization: be correct only with high probability

Streaming model: sketches

- Data is assumed to come as a stream of values
 - e.g. bytes seen when reading off a tape-drive
 - destination IPs seen by a network switch
- Size of universe/stream is much large compared to available memory
 - typically assume memory is $\text{poly}(\log)$
 - Can make limited (possibly single) pass over data
 - Will create a “sketch” : a summary data structure used to answer queries at the end



Streaming problem: distinct count

- Universe is U , number of distinct elements = n , stream size is m
 - Example: U = all IP addresses

10.1.21.10, 10.93.28.1,.....,98.0.3.1,.....10.93.28.1.....

- IPs can repeat
- Want to estimate the number of distinct elements in the stream



Other applications

- Universe = set of all k-grams, stream is generated by document corpus
 - need number of distinct k-grams seen in corpus
- Universe = telephone call records, stream generated by tuples (caller, callee)
 - need number of phones that made > 0 calls



Solutions

- Naïve solution : $O(n \log(U))$ space
 - store all the elements, sort and count distinct
 - store a hash map, insert only if not present in map
- Bit array: $O(|U|)$ space
 - bits initialized to 1 only if element seen in stream
- Can we do this in less space? Not when exact solution needed!!

Approximations

- (ϵ, δ) –approximations
 - Algorithm will use random hash functions
 - Will return an answer \hat{n} such that

$$(1 - \epsilon)n \leq \hat{n} \leq (1 + \epsilon)n$$

- This will happen with probability $1 - \delta$ over the randomness of the algorithm

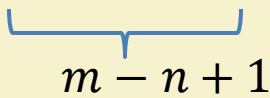


First effort

- Stream length: m , universe size: n
- Proposed algo: Given space S , sample S items from the stream
 - Find the number of distinct elements in this set: \hat{n}
 - return $\hat{n} \times \frac{m}{S}$



First effort

- Stream length: m , distinct elements: n
- Proposed algo: Given space S , sample S items from the stream
 - Find the number of distinct elements in this set: \hat{n}
 - return $\hat{n} \times \frac{m}{S}$
- Not a constant factor approximation
 - $1, 1, 1, 1, \dots, 1, 2, 3, 4, \dots, n-1$

$$m - n + 1$$

Linear Counting

- Bit array B of size m , initialized to all zero
- Hash function $h: [u] \longrightarrow [m]$
- When seeing item x , set $B[h(x)] = 1$

NPTEL



Linear Counting

- Bit array B of size m , initialized to all zero
- Hash function $h: [n] \rightarrow [m]$
- When seeing item x , set $B[h(x)] = 1$
- $Z_m =$ Number of zero entries
- Return estimate $-m \log\left(\frac{Z_m}{m}\right)$

Linear Counting Analysis

- $\Pr[\text{position remaining } 0] = \left(1 - \frac{1}{m}\right)^n \approx e^{-\frac{n}{m}}$
- Expected number of positions at zero = $E[z_m] = me^{-n/m}$
- Using tail inequalities we can show this is concentrated
- Typically useful only for $m = \Theta(n)$, often useful in practice



Flajolet Martin Sketch

- Components
 - “random” hash function $h: U \rightarrow 2^\ell$ for some large ℓ
 - $h(x)$ is a ℓ –length bit string
 - initially assume it is completely random, can relax
- $zero(v)$ = position of rightmost 1 in bit representation of v
= $\max\{ i, 2^i \text{ divides } v \}$
 - $zeros(10110) = 1$, $zeros(110101000) = 3$

Flajolet Martin Sketch

Initialize:

- Choose a “random” hash function $h: U \rightarrow 2^\ell$
- $z \leftarrow 0$

Process(x)





- if $\text{zeros}(h(x)) > z$, $z \leftarrow \text{zeros}(h(x))$

Estimate:

- return $2^{z+1/2}$

Example



	$h(.)$
	0110101
	1011010
	1000100
	1111010

NPTEL



Space usage

- We need $\ell \geq C \log(n)$ for some $C \geq 3$, say
 - by birthday paradox analysis, no collisions with high prob
- Sketch : z , needs to have only $O(\log \log n)$ bits !!!
- Total space usage = $O(\log n + \log \log n)$

Intuition

- Assume hash values are uniformly distributed
- The probability that a uniform bit-string
 - is divisible by 2 is $\frac{1}{2}$
 - is divisible by 4 is $\frac{1}{4}$
 -
 - is divisible by 2^k is $\frac{1}{2^k}$
- We don't expect any of them to be divisible by $2^{\log_2(n)+1}$

Formalizing intuition

- S = set of elements that appeared in stream
- For any $r \in [\ell], j \in S$, X_{rj} = indicator of $\text{zeros}(h(j)) \geq r$
- Y_r = number of $j \in S$ such that $\text{zeros}(h(j)) \geq r$

$$Y_r = \sum_{j \in S} X_{rj}$$

- Let \hat{z} be final value of z after algo has seen all data

Proof of FM

- $Y_r > 0 \iff \hat{z} \geq r$, equivalently, $Y_r = 0 \iff \hat{z} < r$

NPTEL



Proof of FM

- $Y_r > 0 \iff \hat{z} \geq r$, equivalently, $Y_r = 0 \iff \hat{z} < r$

- $E[Y_r] = \sum_{j \in S} E[X_{rj}]$ $X_{rj} = \begin{cases} 1 & \text{with prob } \frac{1}{2^r} \\ 0 & \text{else} \end{cases}$

- $E[Y_r] = \frac{n}{2^r}$ $\text{var}(Y_r) = \sum_{j \in S} \text{var}(X_{rj}) \leq \sum_{j \in S} E[X_{rj}^2]$

Proof of FM

- $\text{var}(Y_r) \leq \sum_{j \in S} E[X_{rj}^2] \leq n/2^r$

$$\Pr[Y_r > 0] = \Pr[Y_r \geq 1] \leq \frac{E[Y_r]}{1} = \frac{n}{2^r}$$

NPTEL



Proof of FM

- $\text{var}(Y_r) \leq \sum_{j \in S} E[X_{rj}^2] \leq n/2^r$

$$\Pr[Y_r > 0] = \Pr[Y_r \geq 1] \leq \frac{E[Y_r]}{1} = \frac{n}{2^r}$$

$$\Pr[Y_r = 0] \leq \Pr[|Y_r - E[Y_r]| \geq E[Y_r]] \leq \frac{\text{var}(Y_r)}{E[Y_r]^2} \leq \frac{2^r}{n}$$



Upper bound

Returned estimate $\hat{n} = 2^{\hat{z}+1/2}$

a = smallest integer with $2^{a+1/2} \geq 4n$

$$\Pr[\hat{n} \geq 4n] = \Pr[\hat{z} \geq a] = \Pr[Y_a > 0] \leq \frac{n}{2^a} \leq \frac{\sqrt{2}}{4}$$



Lower bound

Returned estimate $\hat{n} = 2^{\hat{z}+1/2}$

b = largest integer with $2^{b+1/2} \leq n/4$

$$\Pr \left[\hat{n} \leq \frac{n}{4} \right] = \Pr[\hat{z} \leq b] = \Pr[Y_{b+1} = 0] \leq \frac{2^{b+1}}{n} \leq \frac{\sqrt{2}}{4}$$



Understanding the bound

- By union bound, with prob $1 - \frac{\sqrt{2}}{2}$,

$$\frac{n}{4} \leq \hat{n} \leq 4n$$

- Can get somewhat better constants
- Need only 2-wise independent hash functions, since we only used variances

Improving the probabilities

- To improve the probabilities, a common trick: **median of estimates**
- Create $\hat{z}_1, \hat{z}_2, \dots, \hat{z}_k$ in parallel
 - return median
- Expect at most $\frac{\sqrt{2}}{4}$ of them to exceed $4n$

Improving the probabilities

- To improve the probabilities, a common trick: **median of estimates**
- Create $\hat{z}_1, \hat{z}_2, \dots, \hat{z}_k$ in parallel
 - return median
- Expect at most $\frac{\sqrt{2}}{4} k$ of them to exceed $4n$
- But if median exceeds $4n$, then $\frac{k}{2}$ of them does \rightarrow using Chernoff bound this prob is $\exp(-\Omega(k))$

Improving the probabilities

- To improve the probabilities, a common trick: **median of estimates**
- Create $\hat{z}_1, \hat{z}_2, \dots, \hat{z}_k$ in parallel
 - return median
- Using Chernoff bound, can show that median will lie in $\left[\frac{n}{4}, 4n\right]$ with probability $1 - \exp(-\Omega(k))$.
- Given error prob δ , choose $k = O(\log(\frac{1}{\delta}))$

Summary

- Streaming model– useful abstraction
 - Estimating basic statistics also nontrivial
- Estimating number of distinct elements
 - Linear counting
 - Flajolet Martin



k-MV sketch

- Developed in an effort to get better accuracy
- Additional capabilities for estimating cardinalities of union and intersection of streams
 - If S_1 and S_2 are two streams, can compute their union sketch from individual sketches of S_1 and S_2

[kMV sketch slides courtesy Cohen-Wang]





Sampling via hashing: Thought experiment

- Suppose $h: U \rightarrow [0,1]$ is random hash function such that $h(x) \sim U[0,1]$ for all $x \in U$
- Maintain min-hash value y
 - initialize $y \leftarrow 1$
 - For each item x_i , $y \leftarrow \min(y, h(x_i))$

[kMV sketch slides courtesy Cohen-Wang]

Example



	$h(.)$
	0110101
	1011010
	1000100
	1111010

NPTEL



Intuition

- What information does y have about the number of distinct elements n ?
- Expectation of minimum is $E[\min_i h(x_i)] = \frac{1}{n+1}$



Why is expectation of min = $\frac{1}{n+1}$?

- Imagine a circle instead of $[0, 1]$
- Choose $n + 1$ points uniformly at random

NPTEL



Why is expectation of min = $\frac{1}{n+1}$?

- Imagine a circle instead of $[0, 1]$
- Choose $n + 1$ points uniformly at random
- $n + 1$ intervals are formed
- Expected length of each interval is $\frac{1}{n+1}$

Why is expectation of min = $\frac{1}{n+1}$?

- Imagine a circle instead of $[0, 1]$
- Choose $n + 1$ points uniformly at random
- $n + 1$ intervals are formed
- Expected length of each interval is $\frac{1}{n+1}$
- Think of the first point as the place to cut the circle!

[kMV sketch slides courtesy Cohen-Wang]

k-minimum value sketch

Initialize:

- $y_1, \dots, y_k \leftarrow 1, \dots, 1$

Process(x):

- For all $j \in [k]$, $y_j \leftarrow \min(y_j, h(x_i))$

Estimate:

- return median-of-means($\frac{1}{y_1}, \dots, \frac{1}{y_k}$)



Median-of-means

- Given (ϵ, δ) , choose $k = \frac{c}{\epsilon^2} \log(\frac{1}{\delta})$
- Group t_1, \dots, t_k into $\log(\frac{1}{\delta})$ groups of size $\frac{c}{\epsilon^2}$ each
- Find $\text{mean}(t_i)$ for each group: $Z_1, \dots, Z_{\log(\frac{1}{\delta})}$
- Return $\hat{n} = \text{median of } Z_1, \dots, Z_{\log(\frac{1}{\delta})}$



Example



	h1	h2	h3	h4
	.45	.19	.10	.92
	.35	.51	.71	.20
	.21	.07	.93	.18
	.14	.70	.50	.25

NPTEL



Complexity

- Total space required = $O(k \log n) = O(\frac{1}{\epsilon^2} \log n \log(\frac{1}{\delta}))$
 - can be improved
 - don't need floating points, can use $h: U \rightarrow 2^\ell$ as before
 - can do with k-wise universal hash functions
- Update time per item = $O(k)$
 - However, can show that most items will not result in updates



Theoretical Guarantees

With probability $1 - \delta$, returns \hat{n} satisfies

$$(1 - \epsilon)n \leq \hat{n} \leq (1 + \epsilon)n$$

Proof is simple application of expectation and Chernoff bound



Merging

- For two stream S_1 and S_2 use same set of hash functions
- For each $j \in [k]$, find $\min(y_j, y'_j)$
- Gives estimate of $|S_1 \cup S_2|$



References:

- Primary reference for this lecture
 - Lecture notes by Amit Chakrabarti: <http://www.cs.dartmouth.edu/~ac/Teach/data-streams-lecnotes.pdf>
- Others
 - Blum, Hopcroft, Kannan.
 - Sketch techniques for approximate query processing, Graham Cormode.
<http://dimacs.rutgers.edu/~graham/pubs/papers/sk.pdf>

Thank You!!





IIT KHARAGPUR
IIT GANDHINAGAR



NPTEL ONLINE
CERTIFICATION COURSES

Scalable Data Science

Lecture 9: Frequent Elements

Anirban Dasgupta

Computer Science and Engineering

IIT GANDHINAGAR



IIT Gandhinagar

Indian Institute of
Technology Gandhinagar

Streaming model revisited

- Data is seen as incoming sequence
 - can be just element-ids, or ids +frequency updates
- Arrival only streams
- Arrival + departure
 - Negative updates to frequencies possible
 - Can represent fluctuating quantities, e.g.



Frequency Estimation

- Given the input stream, answer queries about item frequencies at the end
 - Useful in many practical applications e.g. finding most popular pages from website logs, detecting DoD attacks, database optimization



- Also used as subroutine in many problems
 - Entropy estimation, itemset mining etc
- [Slides courtesy of Graham Cormode]

Frequency estimation

Q1. Can we create a data structure, sketch, sublinear in the data size to answer all frequency queries accurately?

NPTEL



Frequency estimation in one pass

Q1. Can we create a data structure, sketch, sublinear in the data size to answer all frequency queries accurately?

– No

Q2. Can we create a sketch to estimate frequencies of the “most frequent” elements exactly?

NPTEL



Frequency estimation in one pass

Q1. Can we create a data structure, sketch, sublinear in the data size to answer all frequency queries exactly?

– No

Q2. Can we create a sketch to answer frequencies of the “most frequent” elements exactly?

– No

Q3. Sketch to estimate frequencies of “most frequent” elements approximately?

Frequency estimation in one pass

Q1. Can we create a data structure, sketch, sublinear in the data size to answer all frequency queries exactly?

- No

Q2. Can we create a sketch to answer frequencies of the “most frequent” elements exactly?

- No

Q3. Sketch to estimate frequencies of “most frequent” elements approximately?

- YES!

Approximate Heavy Hitters

- Given an update stream of length m , find out all elements that occur “frequently”
 - e.g. at least 1% of the time
 - cannot be done in sublinear space, one pass
- Find out elements that occur at least ϕm times, and none that appears $< (\phi - \epsilon)m$ times
 - Error ϵ
 - Related question: estimate each frequency with error $\pm \epsilon m$



Starting with a puzzle

[J. Algorithms, 1981] Suppose we have a list of N numbers, representing votes of N processors on result of some computation. We wish to decide if there is a majority vote and what that vote is.

- By J.S. Moore
- Did not talk about streaming solution, but proposed solution is
- Strict majority: $>N/2$



Majority Algorithm

- Arrivals only model
- Start with a counter set to zero
- For each item
 - if counter = 0, pick new item and increment counter
 - else if new item is same as item in hand, increment counter
 - else decrement counter



Majority Algorithm

- Start with a counter set to zero
- For each item
 - if counter = 0, pick new item and increment counter
 - else if new item is same as item in hand, increment counter
 - else decrement counter
- If there is a majority item, it is in hand at the end
- Proof: Since majority occurs $> N/2$ times, not all occurrences can be cancelled out



Frequent [Misra-Gries]

- Keep k counters and items in hand

Initialize:

- Set all counters to 0

Process(x)

- if x is same as any item in hand, increment its counter
- else if number of items $< k$, store x with counter = 1
- else drop x and decrement all counters

Query(q)

- If q is in hand return its counter, else 0

Frequent

- f_x be the true frequency of element x
- At the end, some set of elements is stored with counter values
- If *query* y in hand, $\hat{f}_y = \text{counter value}$, else $\hat{f}_y = 0$



Example



NPTEL



Theoretical Bound

Claim: No element with frequency $> m/k$ is missed at the end

NPTEL



Theoretical Bound

Claim: No element with frequency $> m/k$ is missed at the end

Intuition: Each decrement (including drop) is charged with k arrivals. Therefore, will have some copy of an item with frequency $> m/k$



Stronger Claim

Choose $k = \frac{1}{\epsilon}$. For every item x , with frequency f_x the algo can return an estimate \hat{f}_x such that

$$f_x - \epsilon m \leq \hat{f}_x \leq f_x$$



Stronger Claim

Choose $k = \frac{1}{\epsilon}$. For every item x , with frequency f_x the algo can return an estimate \hat{f}_x such that

$$f_x - \epsilon m \leq \hat{f}_x \leq f_x$$

Same intuition, whenever we drop a copy of item x , we also drop $k - 1$ copies of other items



Summary

- Simple deterministic algorithm to estimate heavy hitters
 - Works only in the arrival model
- Proposed in 1982, rediscovered multiple times with modifications
- Also basis of matrix low rank approximation
- Our next lecture will discuss other algorithms



References:

- Primary references for this lecture
 - Lecture slides by Graham Cormode
<http://dmac.rutgers.edu/Workshops/WGUnifyingTheory/Slides/cormode.pdf>
 - Lecture notes by Amit Chakrabarti: <http://www.cs.dartmouth.edu/~ac/Teach/data-streams-lecnotes.pdf>
 - Sketch techniques for approximate query processing, Graham Cormode.
<http://dimacs.rutgers.edu/~graham/pubs/papers/sk.pdf>



Thank You!!



IIT Gandhinagar
Indian Institute of
Technology Gandhinagar



NPTEL ONLINE
CERTIFICATION COURSES

Anirban Dasgupta
Computer Science and Engg.